

EDUCATION

The Chinese University of Hong Kong, Shenzhen · B.S. in Data Science and Big Data Technology Sep 2023 - Jun 2027

Relevant coursework: Data Structures, Database Systems, Machine Learning, Optimization, Stochastic Processes, Probability & Statistics, Linear Algebra and Discrete Mathematics. English-medium academic environment with technical reading and writing experience.

TECHNICAL SKILLS

Applied AI	RAG / hybrid retrieval, tool calling, agent workflows, human-in-the-loop, evaluation, tracing / observability
Engineering	Python, FastAPI, REST APIs, SQLite, Git; reproducible demos, tests and API-layer implementation
Data & ML	NumPy, Pandas, Scikit-learn, LightGBM; data cleaning, feature engineering, model evaluation and time-series validation
Other	C/C++, SQL, Jupyter, VS Code, LaTeX; OpenAI API, Claude Code and Antigravity for AI-assisted engineering workflows

SELECTED AI PROJECTS

Evidence Research Copilot · Independent Project 2026

- RAG · BM25 · TF-IDF · Reciprocal Rank Fusion · Evaluation · FastAPI
- Built an evidence-grounded research assistant that chunks local documents, retrieves with BM25 and TF-IDF, and fuses independent rankings with Reciprocal Rank Fusion.
 - Exposed source chunks, ranking signals and citations through the answer pipeline so retrieval evidence remains inspectable instead of hidden behind a black-box response.
 - Added a reproducible regression evaluation suite for Recall@K, MRR and citation validity, plus a FastAPI/browser demo that runs without a paid model API.

Agent Workflow Studio · Independent Project 2026

- AI Agents · Tool Calling · Approval Gates · Retry / Budget · SQLite Tracing · FastAPI
- Implemented a framework-light tool-using agent runtime with a planner, typed tool registry, execution budget, retry policy and explicit workflow state.
 - Added risk classes and human approval gates before side-effect-class tools; used a restricted AST evaluator for arithmetic rather than unsafe eval.
 - Persisted inputs, outputs, latency, errors and status to SQLite traces and surfaced them in a web execution timeline; tests cover tool execution and approval pause/continue behavior.

Intelligent Email Auto-Reply System · Core Developer 2025

- LLM Application · NLP · Confidence Routing · Human Review
- Designed an LLM-based customer-service email workflow for intent recognition, fine-grained classification and reply generation, with confidence-based routing between automation and human review.
 - Implemented the model capability inside a runnable frontend/backend system, connecting model calls with business routing and a human fallback path.

COMPETITION

Mathematical Contest in Modeling / Interdisciplinary Contest in Modeling · Successful Participant 2026

Contributed to mathematical modeling, data analysis, parameter solving, visualization and an English technical paper under strict time constraints.